

Exploring the use of aligned and disaligned multimodal features in MultiModal Recommender Systems

Fontana Emanuele

October 29, 2025

Contents

1	Introduction	3
1.1	Recommender Systems	3
1.2	Multimodal Recommender Systems	4
2	Aligned multimodal models	6
2.1	Contrastive Learning and AudioCLIP	6
2.1.1	Contrastive Learning	6
2.1.2	CLIP	6
2.1.3	AudioCLIP	7
2.2	Contrastive Learning arcitetures in Multimodal Recommender Systems	8
2.2.1	CLIP in Multimodal Recommender Systems	8
2.2.2	AudioCLIP in Multimodal Recommender Systems	10
3	Experimental Settings	11
3.1	Research Questions	11
3.2	Datasets	11
3.2.1	MovieLens 1M	11
3.2.2	LastFM	11
3.3	Multimodal Embedding Extraction Pipeline	12
3.3.1	Overview	12
3.3.2	File Organization and Naming Convention	12
3.3.3	Embedding Models	13
3.3.4	Extraction Process	13

3.3.5	Output Files	14
3.4	Interaction Filtering and Remapping Pipeline	14
3.4.1	Common Pipeline Steps	14
3.4.2	MovieLens Mapping	15
3.4.3	LastFM Mapping	16
4	Experiments	18
4.1	MovieLens_1M	18
4.2	LastFM	19

1 Introduction

1.1 Recommender Systems

With **Recommender Systems** we refer to all those techniques that have the effect of guiding users in personalized way to interesting objects in a large space of possible options[5]. There are different types of Recommender Systems:

- **Content-Based Filtering:** This approach recommends items similar to those the user has liked in the past. It uses item features and user preferences to make recommendations.
- **Collaborative Filtering:** This method recommends items based on the preferences of similar users. It can be user-based or item-based.
 - **User-Based Collaborative Filtering:** This method finds users with similar preferences and recommends items they have liked.
 - **Item-Based Collaborative Filtering:** This method finds items similar to those the user has liked and recommends them. With respect to Content-Based Filtering, this method doesn't use item features to make recommendations.
- **Hybrid Methods:** These methods combine multiple recommendation techniques to improve accuracy and overcome the limitations of individual approaches.
- **Knowledge-Based Recommender Systems:** These systems use explicit knowledge about users and items to make recommendations.
- **Context-Aware Recommender Systems:** These systems take into account contextual information, such as time, location, or social context, to provide more relevant recommendations.




Figure 1: Conceptual representation of a Recommender Systems

Recommender Systems are not perfect. They can suffer of several problems like:

- **Cold Start Problem:** Because new users often have very few records in the system, it is hard to guess their preferences given the insufficient information[8]
- **Data Sparsity:** Most users use the system but do not give rating for feedback to the system in a proper way.[6]. In addition, in many real-world application we have thousand if not milion of user and item, so it's impossible to think that all user with interact with all item.
- **Scalability:** As the number of users and items increases, the computational complexity of recommendation algorithms can become a challenge. So we can divide scalability problems in two categories:
 - **Software Scalability:** There's a need to develop algorithms that can handle milions of users and items
 - **Hardware Scalability:** There's a need to create systems with powerful hardware that can handle the computational cost in time and space
- **Over Specialization :** Sometimes Recommender algorithms can "over-fit" users' preferences, which lead to suggest items that are too similar to those already liked by the user not allowing him/her to discover new items that might be relevant

1.2 Multimodal Recommender Systems

Differently from classical recommendation approaches, multimodal recommendation exploits multimodal side information about items to enrich the interaction matrix, alleviating its sparsity, and to better understand the content to be recommended; these advantages, in general, lead to more accurate and precise recommendations[9]. We can think about different types of multimodal information: **Text** in natural language that can be used in every context, **Audio** for Music, **Images** for products such as clothes or furnitures, **Videos** for movies and so on.

The powerfulness of multimodal recommendation is given by the fact that different modalities can provide information about the same item from different perspectives. For example, in a movie recommender system, the textual description of a movie can provide information about its plot, screenshots

from the film can provide visual clues about the style of photography and the type of film (animated, live-action, stop-motion, etc.)

Some DNN architectures known as **Encoders** are used to extract features from different modalities that can be used in different ways like:

- **Concat:** Refers to the concatenation of multimodal features. [1]
- **Fusion and attention:** Refers to the attention mechanisms (like self-attention) to combine multimodal features. [4]
- **GNNs:** They use Graph Neural Networks to model the relationships between items and users, incorporating multimodal features into the graph structure




Figure 2: Conceptual representation of a Multimodal Recommender Systems

2 Aligned multimodal models

2.1 Contrastive Learning and AudioCLIP

2.1.1 Contrastive Learning

Contrastive Learning is self-supervised representation learning by training a model to differentiate between similar and dissimilar samples.[3]. In few words, contrastive learning is a technique used to learn a feature space where similar samples are close together and dissimilar samples are far apart. To achieve this, a specific loss function known as **Contrastive Loss** is used:

$$\ell_{i,j} = -\log \frac{\exp(\text{sim}(z_i, z_j)/\tau)}{\sum_{k=1}^{2N} [k \neq i] \exp(\text{sim}(z_i, z_k)/\tau)}, \quad (1)$$

where $\text{sim}(z_i, z_j)$ represents the similarity between two feature vectors z_i and z_j , τ is a temperature parameter, and N is the batch size. The numerator focuses on the similarity between the positive pair (z_i, z_j) , while the denominator normalizes over all pairs excluding z_i itself.




Figure 3: Conceptual representation of Contrastive Learning

2.1.2 CLIP

CLIP, which stands for Contrastive Language-Image Pre-training, is a neural network architecture developed by OpenAI[7] that uses Contrastive Learning

to connect textual and visual information. The main idea behind CLIP is that visual and textual embeddings live in the same feature space. Given the textual description of an image and the image itself, the embeddings returned by the model should be close together in the feature space. In few words, CLIP is trained to predict which images are most relevant to a given text prompt, and vice versa.




Figure 4: Conceptual representation of CLIP architecture

2.1.3 AudioCLIP

AudioCLIP [2] is an extension of the CLIP model that incorporates an additional audio modality, enabling it to understand and relate audio data alongside text and images. The architecture of AudioCLIP consists of three main components:

- **A CLIP Model** to handle text and image modalities. In particular a ResNet based CLIP model is used
- **Audio Encoder** to handle audio modality. In particular an ESRes-NetXt model is used

The idea to realize AudioCLIP is very simple: in addition to the text-image contrastive loss used in CLIP, two additional loss were added: **text-audio** and **image-audio** contrastive loss. In this way the three modalities are connected in the same feature space. The training procedure of AudioCLIP consists of different main steps:

- **Audio-Head Pre-Training**
 - **Standalone:** The audio head (based on ESResNeXt) is first pre-trained independently on the *AudioSet* dataset.

- **Cooperative:**
 - * The classification layer of the pre-trained audio head is replaced with a randomly initialized layer whose output size matches CLIP’s embedding space.
 - * The audio head is then trained jointly with the frozen text and image heads, in a *multi-modal knowledge distillation* setup, making audio embeddings compatible with CLIP embeddings.

- **AudioCLIP Training**

- After making the audio head compatible with CLIP, the entire tri-modal model (audio, text, image) is trained on *AudioSet*.
- All three modality-specific heads are updated together, allowing the model to adapt to the distributions of audio samples, image frames, and textual class names.
- This joint training improves performance compared to training only the audio head.




Figure 5: Conceptual representation of AudioCLIP architecture

2.2 Contrastive Learning architectures in Multimodal Recommender Systems

CLIP and AudioCLIP were born, and mainly used, for classification tasks. However, their ability to create a shared feature space for different modalities makes them suitable for multimodal recommender systems as well.

2.2.1 CLIP in Multimodal Recommender Systems

Zixuan Yi et al. [10] analyzed the use of CLIP in multimodal recommender systems. Up to few years ago, multimodal features were extracted using pretrained models on a single modality, such as ResNet for images and BERT

for text. This leads to a problem: the features extracted from different modalities are not aligned in the same feature space. To solve this problem, they proposed to use CLIP to extract aligned image and text features. In particular, they used both a frozen and a fine-tuned CLIP model to extract image and text features. Experiments involved five multimodal recommender systems: VBPR, MMGCN, MMGCL, SLMRec and LATTICE and showed that both frozen and fine-tuned CLIP features significantly outperformed the traditional pretrained features in four out of five models. To be more specific, due to its learning objective, LATTICE performs better without CLIP features. The fine-tune algorithm proposed by the authors is known as **END-TO-END Training**.

End-To-End Training The end-to-end fine-tuning procedure integrates the CLIP encoder directly into the recommendation model and jointly optimizes both components using recommendation losses. The training steps are as follows:

1. **Load Data:** Prepare the dataset by loading raw data
2. **Initialize CLIP Encoder:** Load a pre-trained CLIP model and its corresponding weights.
3. **Generate Embeddings:** Use the CLIP encoder to extract image embeddings and text embeddings from the dataset:
4. **Integrate Embeddings:** Feed the extracted embeddings into the recommendation model as initial item representations.
5. **Joint Optimisation:** For each training epoch, jointly update both the CLIP encoder and the recommendation model:
 - (a) Perform a forward pass to compute user-item scores.
 - (b) Compute the recommendation loss
 - (c) Backpropagate the loss and update both the CLIP encoder and recommendation model parameters:
6. **Evaluation:** After each epoch, evaluate and log recommendation performance metrics (e.g., NDCG, Recall) to monitor progress.

This end-to-end strategy allows the CLIP encoder to adapt its visual and textual representations specifically to the recommendation task, leading to improved alignment between modalities and better overall recommendation performance.

2.2.2 AudioCLIP in Multimodal Recommender Systems

In literature there are no evidence of the use of AudioCLIP in multimodal recommender systems. Also, there no evidence of the use of CLIP + disaligned audio features in multimodal recommender systems

3 Experimental Settings

3.1 Research Questions

The main research question that this project aims to address is:

- **RQ1:** Can the integration of aligned multimodal embeddings (text, image, audio) enhance the performance of recommender systems compared to unimodal or non-aligned multimodal approaches?

3.2 Datasets

This project employs two benchmark datasets for evaluating multimodal recommendation:

3.2.1 MovieLens 1M

MovieLens 1M is a widely-used collection of explicit movie ratings containing **User-movie interactions**: implicit feedback from user ratings. For each movie, the dataset has been enriched with multimodal data:

- **Images:** movie posters (.jpg format)
- **Audio:** movie trailer (.wav format)
- **Text:** textual descriptions, synopses, or reviews (.txt format)

3.2.2 LastFM

LastFM is a music listening dataset containing:

- **User-artist interactions:** implicit feedback from listening history
- **Multiple tracks per artist:** track-level multimodal features

For each track, the dataset includes:

- **Images:** album artwork (.jpg format)
- **Audio:** songs (.wav format)
- **Text:** track metadata or descriptions (.txt format)

Note: LastFM presents a unique challenge where interactions are recorded at the artist level, while multimodal features are extracted at the track level. This mismatch is addressed in the mapping pipeline (Section 3.4).

3.3 Multimodal Embedding Extraction Pipeline

The embedding extraction pipeline is **dataset-agnostic** and applies the same process to both MovieLens and LastFM. This section describes the unified extraction methodology.

3.3.1 Overview

The extraction pipeline performs the following steps:

1. File discovery and alignment across modalities
2. Validation and corruption detection
3. Per-modality embedding extraction using multiple encoder architectures
4. Strict consistency enforcement across modalities
5. Output generation in aligned NumPy arrays

3.3.2 File Organization and Naming Convention

All multimodal files are organized with a consistent basename structure. For both datasets:

- Each item (movie/track) is identified by a unique ID
- Files across modalities share the same basename

Example (MovieLens): For movie ID 1:

- Image: `ml1m/_images/1.jpg`
- Audio: `ml1m/_audios/1.wav`
- Text: `ml1m/_texts/1.txt`

Example (LastFM): For artist 1000 with track variant 1:

- Image: `lastfm/_images/1000_1.jpg`
- Audio: `lastfm/_audios/1000_1.wav`
- Text: `lastfm/_texts/1000_1.txt`

(there are up to 5 songs per artist)

3.3.3 Embedding Models

Multiple encoder architectures are employed to extract embeddings, providing both aligned and specialized representations:

Aligned Multimodal Models

- **AudioCLIP**: Extracts aligned embeddings across all three modalities (image, audio, text) in a shared feature space
- **CLIP**: Extracts aligned image-text embeddings

Modality-Specific Models

- **ViT (Vision Transformer)**: Image embeddings
- **VGGish**: Audio embeddings
- **MiniLM**: Text embeddings

3.3.4 Extraction Process

Step 1: File Discovery and Alignment The script scans three input directories containing raw multimodal data:

- Identifies files by basename (e.g., movie ID or artist_track combination)
- Builds maps from basename to absolute file paths
- Computes the **intersection** of basenames across all three modalities

Only items present in all three modalities are considered candidates for extraction. This ensures row-wise correspondence in the final embedding arrays.

Step 2: Audio Validation Before extraction, audio files undergo validation to detect corrupted or unreadable files. Invalid audio files are logged and excluded from processing. Exclusion is global: if an audio file fails validation, the corresponding image and text files are also removed to maintain alignment.

Step 3: Extraction Files that pass validation are processed in order to extract embeddings. If the extraction for a specific modality fails (e.g., due to file corruption), the entire item is excluded to ensure strict alignment.

3.3.5 Output Files

The extraction pipeline produces the following files:

Embedding Arrays (NumPy format) Per-modality embeddings for each encoder:

- AudioCLIP: `image_audioclip.npy`, `audio_audioclip.npy`, `text_audioclip.npy`
- CLIP: `image_clip.npy`, `text_clip.npy`
- ViT: `image_vit.npy`
- VGGish: `audio_vggish.npy`
- MiniLM: `text_minilm.npy`

Metadata

- `item_features.csv`: Maps item identifier (basename) to row index in embedding arrays
- Format: two columns (`item_id`, `row_index`)
- **Critical property**: Row i in each `.npy` file corresponds to the same item

3.4 Interaction Filtering and Remapping Pipeline

After embedding extraction, the interaction datasets must be filtered and remapped to ensure consistency with the available multimodal features. This section describes the **dataset-specific** mapping procedures.

3.4.1 Common Pipeline Steps

Both datasets follow the same core logic:

1. Load original interaction file and `item_features.csv`
2. Filter interactions to include only items with complete multimodal embeddings
3. Apply 5-core filtering to ensure interaction density
4. Remap user and item IDs to contiguous ranges $[0, \dots, n-1]$

5. Reconstruct embedding arrays in the new item order
6. Generate output files with remapped IDs and embeddings

The key difference lies in **Step 2**, where LastFM requires special handling for artist-to-track mapping.

3.4.2 MovieLens Mapping

Step 1: Loading Data The script loads:

- `movielens_1m.inter`: Original interaction file (user_id, item_id, rating, timestamp)
- `item_features.csv`: Mapping from movie ID to embedding row index
- All `.npz` embedding files

Step 2: Feature-Based Filtering Check between `movielens_1m.inter` and `item_features.csv` to retain only interactions referencing movies contained in both files.

Step 3: 5-Core Filtering To address data sparsity, the script applies iterative 5-core filtering:

Algorithm:

1. Count interactions per user and per item
2. Remove users with < 5 interactions
3. Remove items with < 5 interactions
4. Repeat until convergence (no further removals)

This ensures that every user has rated at least 5 items and every item has been rated by at least 5 users.

Step 4: ID Remapping After filtering, IDs may be non-contiguous. The script creates contiguous mappings:

- **Users:** $\{u_1, u_2, \dots, u_m\} \rightarrow \{0, 1, \dots, m - 1\}$
- **Items:** $\{i_1, i_2, \dots, i_n\} \rightarrow \{0, 1, \dots, n - 1\}$

Step 5: Embedding Reindexing The original embedding arrays are indexed according to `item_features.csv`. After remapping, embeddings must be reordered to match the new item indices:

Algorithm:

1. For each new item index $j \in [0, n - 1]$:
 - Lookup original item ID: $i_{\text{old}} = \text{item_mapping}^{-1}[j]$
 - Find row index in original embeddings: $r = \text{item_features}[i_{\text{old}}]$
 - Copy embedding vectors: $\text{embedding_new}[j] = \text{embedding_old}[r]$
2. Save reindexed embeddings

3.4.3 LastFM Mapping

LastFM requires special handling due to a structural mismatch:

- **Interactions** are recorded at the **artist level** (user listens to artist)
- **Embeddings** are extracted at the **track level** (one feature set per track)

Resolution Strategy: Artist-to-Track Expansion The mapper resolves this mismatch by:

1. Mapping each artist to all available track variants in `item_features.csv`
2. Expanding each user-artist interaction into multiple user-track interactions
3. Treating each variant as a separate item in the recommendation task

Step 1: Loading Data The script loads:

- `lastfm.inter`: Original interaction file (user_id, artist_id, listen_count, timestamp)
- `item_features.csv`: Mapping from `artist_track` identifier to embedding row index
- All `.npz` embedding files

Step 2: Artist-to-Track Mapping and Expansion The script performs two sub-steps:

Sub-step 2a: Build Artist-to-Tracks Dictionary

- Parse `item_features.csv` to extract artist and track identifiers
- Group tracks by artist: $\text{artist} \rightarrow [\text{track}_1, \text{track}_2, \dots]$

Sub-step 2b: Expand Interactions For each user-artist interaction :

- Lookup available tracks for artist a : $T_a = [\text{track}_1, \dots, \text{track}_k]$
- Create k new interactions: $(u, \text{track}_1), \dots, (u, \text{track}_k)$

This expansion significantly increases the number of interactions.

Step 3–5: Standard Pipeline After expansion, the pipeline follows the same steps as MovieLens:

- 5-core filtering (now at track level)
- Contiguous ID remapping
- Embedding reindexing

4 Experiments

This section reports test-set results for several models and multimodal configurations. Tables are grouped by cutoff k : one table for @5, @10, @20 and @50. For each row we report Recall, Precision, NDCG and MAP at the corresponding cutoff.

4.1 MovieLens_1M

Feature extractors used Unless explicitly specified in a table row as *CLIP* or *AudioCLIP*, the reported VBPR variants use the following default feature extractors:

- Text: MiniLM (sentence-transformer)
- Image: ViT (Vision Transformer)
- Audio: VGGish-based features

Model	Modality	rec@5	prec@5	ndcg@5	map@5
BPR	None	0.1089	0.1898	0.2096	0.1289
VBPR	Text (MiniLM)	0.1144	0.1999	0.2196	0.1363
VBPR	Audio (VGGish)	0.1147	0.2015	0.2217	0.1374
VBPR	Image (ViT)	0.1148	0.1985	0.2193	0.1356
VBPR	Text+Image (disaligned)	0.1139	0.1972	0.2179	0.1353
VBPR	Text+Image (CLIP)	0.1152	0.1997	0.2205	0.1376
VBPR	Text+Audio+Image (disaligned)	0.1163	0.1986	0.2191	0.1356
VBPR	Text+Audio+Image (AudioCLIP)	0.0940	0.1694	0.1857	0.1125
VBPR	CLIP+audio (disaligned)	0.1898	0.2009	0.2299	0.1384

Table 1: Metrics at cutoff $k = 5$ (test)

Model	Modality	rec@10	prec@10	ndcg@10	map@10
BPR	None	0.1776	0.1631	0.2159	0.1120
VBPR	Text (MiniLM)	0.1877	0.1713	0.2269	0.1192
VBPR	Audio (VGGish)	0.1876	0.1722	0.2278	0.1194
VBPR	Image (ViT)	0.1881	0.1720	0.2280	0.1200
VBPR	Text+Image (disaligned)	0.1895	0.1725	0.2279	0.1199
VBPR	Text+Image (CLIP)	0.1911	0.1745	0.2300	0.1211
VBPR	Text+Audio+Image (disaligned)	0.1876	0.1718	0.2271	0.1195
VBPR	Text+Audio+Image (AudioCLIP)	0.1570	0.1476	0.1922	0.0969
VBPR	CLIP+audio (disaligned)	0.2936	0.1734	0.2562	0.1211

Table 2: Metrics at cutoff $k = 10$ (test)

Model	Modality	rec@20	prec@20	ndcg@20	map@20
BPR	None	0.2743	0.1341	0.2400	0.1112
VBPR	Text (MiniLM)	0.2939	0.1425	0.2544	0.1195
VBPR	Audio (VGGish)	0.2913	0.1432	0.2543	0.1193
VBPR	Image (ViT)	0.2913	0.1424	0.2543	0.1200
VBPR	Text+Image (disaligned)	0.2904	0.1422	0.2534	0.1194
VBPR	Text+Image (CLIP)	0.2942	0.1424	0.2554	0.1204
VBPR	Text+Audio+Image (disaligned)	0.2926	0.1426	0.2542	0.1196
VBPR	Text+Audio+Image (AudioCLIP)	0.2466	0.1225	0.2147	0.0957
VBPR	CLIP+audio (disaligned)	0.4677	0.1432	0.3148	0.1208

Table 3: Metrics at cutoff $k = 20$ (test)

Model	Modality	rec@50	prec@50	ndcg@50	map@50
BPR	None	0.4499	0.0954	0.2991	0.1250
VBPR	Text (MiniLM)	0.4678	0.1001	0.3131	0.1339
VBPR	Audio (VGGish)	0.4675	0.1006	0.3132	0.1335
VBPR	Image (ViT)	0.4654	0.1002	0.3131	0.1344
VBPR	Text+Image (disaligned)	0.4646	0.0999	0.3119	0.1335
VBPR	Text+Image (CLIP)	0.4681	0.1001	0.3141	0.1347
VBPR	Text+Audio+Image (disaligned)	0.4672	0.1001	0.3129	0.1339
VBPR	Text+Audio+Image (AudioCLIP)	0.4176	0.0892	0.2721	0.1082
VBPR	CLIP+audio (disaligned)	0.2223	0.1004	0.2009	0.1351

Table 4: Metrics at cutoff $k = 50$ (test)

4.2 LastFM

Model	Modality	rec@5	prec@5	ndcg@5	map@5
BPR	None	0.0655	0.1577	0.1595	0.1414
VBPR	Text (MiniLM)	0.0733	0.1742	0.1763	0.1585
VBPR	Image (ViT)	0.0754	0.1787	0.1813	0.1611
VBPR	Text+Image (disaligned)	0.0768	0.1831	0.1859	0.1660
VBPR	Text+Image (CLIP)	0.0707	0.1691	0.1714	0.1524
VBPR	Text+Audio+Image (AudioCLIP)	0.0694	0.1711	0.1735	0.1531

Table 5: Metrics at cutoff $k = 5$ (test) for LastFM

Model	Modality	rec@10	prec@10	ndcg@10	map@10
BPR	None	0.1132	0.1404	0.1534	0.1061
VBPR	Text (MiniLM)	0.1268	0.1531	0.1695	0.1200
VBPR	Image (ViT)	0.1298	0.1552	0.1734	0.1214
VBPR	Text+Image (disaligned)	0.1307	0.1560	0.1755	0.1242
VBPR	Text+Image (CLIP)	0.1222	0.1488	0.1646	0.1149
VBPR	Text+Audio+Image (AudioCLIP)	0.1196	0.1482	0.1636	0.1136

Table 6: Metrics at cutoff $k = 10$ (test) for LastFM

Model	Modality	rec@20	prec@20	ndcg@20	map@20
BPR	None	0.1890	0.1161	0.1703	0.0998
VBPR	Text (MiniLM)	0.1990	0.1218	0.1831	0.1119
VBPR	Image (ViT)	0.2086	0.1265	0.1899	0.1142
VBPR	Text+Image (disaligned)	0.2054	0.1243	0.1896	0.1153
VBPR	Text+Image (CLIP)	0.2009	0.1225	0.1820	0.1082
VBPR	Text+Audio+Image (AudioCLIP)	0.1915	0.1189	0.1766	0.1042

Table 7: Metrics at cutoff $k = 20$ (test) for LastFM

Model	Modality	rec@50	prec@50	ndcg@50	map@50
BPR	None	0.3318	0.0814	0.2343	0.1211
VBPR	Text (MiniLM)	0.3494	0.0858	0.2509	0.1344
VBPR	Image (ViT)	0.3566	0.0867	0.2562	0.1367
VBPR	Text+Image (disaligned)	0.3484	0.0853	0.2543	0.1374
VBPR	Text+Image (CLIP)	0.3539	0.0867	0.2508	0.1314
VBPR	Text+Audio+Image (AudioCLIP)	0.3383	0.0835	0.2422	0.1257

Table 8: Metrics at cutoff $k = 50$ (test) for LastFM

References

- [1] Xi Chen, Yangsiyi Lu, Yuehai Wang, and Jianyi Yang. Cmbf: Cross-modal-based fusion recommendation algorithm. *Sensors*, 21(16), 2021.
- [2] Andrey Guzhov, Federico Raue, Jörn Hees, and Andreas Dengel. Audioclip: Extending clip to image, text and audio. In *ICASSP 2022-2022 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, pages 976–980. IEEE, 2022.
- [3] Haigen Hu, Xiaoyuan Wang, Yan Zhang, Qi Chen, and Qiu Guan. A comprehensive survey on contrastive learning. *Neurocomputing*, 610:128645, 2024.
- [4] Peishan Li, Weixiao Zhan, Lutao Gao, Shuran Wang, and Linnan Yang. Multimodal recommendation system based on cross self-attention fusion. *Systems*, 13(1), 2025.
- [5] Pasquale Lops, Marco de Gemmis, and Giovanni Semeraro. *Content-based Recommender Systems: State of the Art and Trends*, pages 73–105. Springer US, Boston, MA, 2011.
- [6] Nitin Mishra, Saumya Chaturvedi, Aanchal Vij, and Sunita Tripathi. Research problems in recommender systems. In *Journal of Physics: Conference Series*, volume 1717, page 012002. IOP publishing, 2021.
- [7] Alec Radford, Jong Wook Kim, Chris Hallacy, Aditya Ramesh, Gabriel Goh, Sandhini Agarwal, Girish Sastry, Amanda Askell, Pamela Mishkin, Jack Clark, et al. Learning transferable visual models from natural language supervision. In *International conference on machine learning*, pages 8748–8763. PmLR, 2021.
- [8] Yu Rong, Xiao Wen, and Hong Cheng. A monte carlo algorithm for cold start recommendation. In *Proceedings of the 23rd international conference on World wide web*, pages 327–336, 2014.
- [9] Giuseppe Spillo, Elio Musacchio, Cataldo Musto, Marco de Gemmis, Pasquale Lops, and Giovanni Semeraro. See the movie, hear the song, read the book: Extending movielens-1m, last.fm-2k, and dbbook with multimodal data. In *Proceedings of the Nineteenth ACM Conference on Recommender Systems, RecSys '25*, page 847–856, New York, NY, USA, 2025. Association for Computing Machinery.

- [10] Zixuan Yi, Zijun Long, Iadh Ounis, Craig Macdonald, and Richard Mccreadie. Large multi-modal encoders for recommendation. *arXiv preprint arXiv:2310.20343*, 2023.